

```
In [1]: import numpy as np
        from sklearn.datasets import fetch_openml
```

```
In [2]: mnist = fetch_openml('mnist_784', version = 1, cache = True, as_frame = False)
        print(mnist.DESCR)
```

```
C:\Users\armen\anaconda3\Lib\site-packages\sklearn\datasets\_openml.py:968: FutureWarning: The default value of `parser` will change from `liac-arff` to `auto` in 1.4. You can set `parser='auto'` to silence this warning. Therefore, an `ImportError` will be raised from 1.4 if the dataset is dense and pandas is not installed. Note that the pandas parser may return different data types. See the Notes Section in fetch_openml's API doc for details.
  warn(
```

****Author**:** Yann LeCun, Corinna Cortes, Christopher J.C. Burges
****Source**:** [MNIST Website](http://yann.lecun.com/exdb/mnist/) - Date unknown
****Please cite**:**

The MNIST database of handwritten digits with 784 features, raw data available at: <http://yann.lecun.com/exdb/mnist/>. It can be split in a training set of the first 60,000 examples, and a test set of 10,000 examples

It is a subset of a larger set available from NIST. The digits have been size-normalized and centered in a fixed-size image. It is a good database for people who want to try learning techniques and pattern recognition methods on real-world data while spending minimal efforts on preprocessing and formatting. The original black and white (bilevel) images from NIST were size normalized to fit in a 20x20 pixel box while preserving their aspect ratio. The resulting images contain grey levels as a result of the anti-aliasing technique used by the normalization algorithm. The images were centered in a 28x28 image by computing the center of mass of the pixels, and translating the image so as to position this point at the center of the 28x28 field.

With some classification methods (particularly template-based methods, such as SVM and K-nearest neighbors), the error rate improves when the digits are centered by bounding box rather than center of mass. If you do this kind of pre-processing, you should report it in your publications. The MNIST database was constructed from NIST's NIST originally designated SD-3 as their training set and SD-1 as their test set. However, SD-3 is much cleaner and easier to recognize than SD-1. The reason for this can be found on the fact that SD-3 was collected among Census Bureau employees, while SD-1 was collected among high-school students. Drawing sensible conclusions from learning experiments requires that the result be independent of the choice of training set and test among the complete set of samples. Therefore it was necessary to build a new database by mixing NIST's datasets.

The MNIST training set is composed of 30,000 patterns from SD-3 and 30,000 patterns from SD-1. Our test set was composed of 5,000 patterns from SD-3 and 5,000 patterns from SD-1. The 60,000 pattern training set contained examples from approximately 250 writers. We made sure that the sets of writers of the training set and test set were disjoint. SD-1 contains 58,527 digit images written by 500 different writers. In contrast to SD-3, where blocks of data from each writer appeared in sequence, the data in SD-1 is scrambled. Writer identities for SD-1 is available and we used this information to unscramble the writers. We then split SD-1 in two: characters written by the first 250 writers went into our new training set. The remaining 250 writers were placed in our test set. Thus we had two sets with nearly 30,000 examples each. The new training set was completed with enough examples from SD-3, starting at pattern # 0, to make a full set of 60,000 training patterns. Similarly, the new test set was completed with SD-3 examples starting at pattern # 35,000 to make a full set with 60,000 test patterns. Only a subset of 10,000 test images (5,000 from SD-1 and 5,000 from SD-3) is available on this site. The full 60,000 sample training set is available.

Downloaded from openml.org.

```
In [3]: #To check data structure first
import matplotlib.pyplot as plt #
import pandas as pd

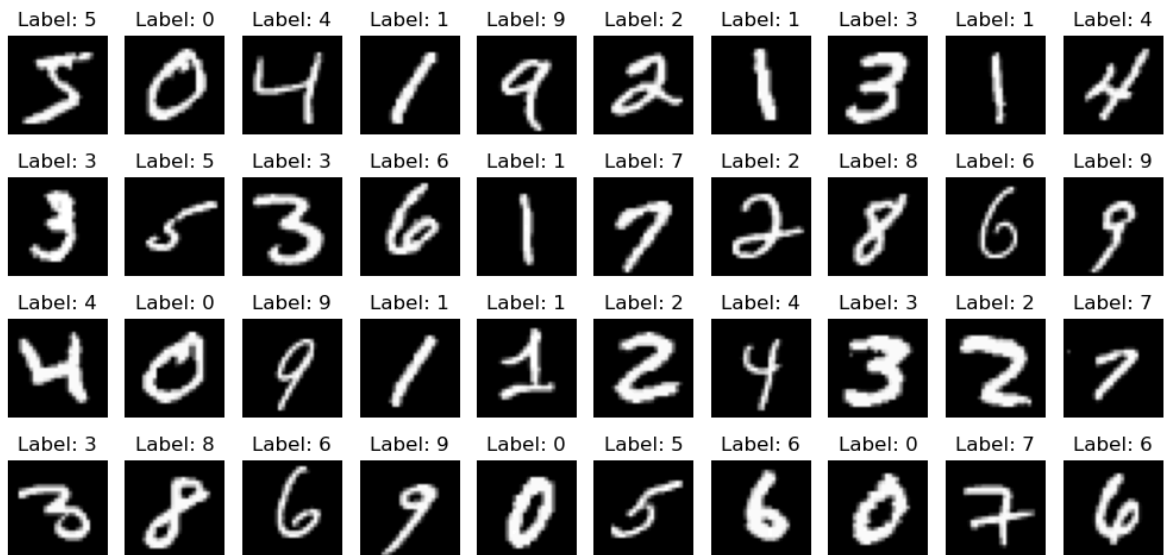
X, y = mnist["data"], mnist["target"]
print("Shape of X:", X.shape)
print("Shape of y:", y.shape)
```

Shape of X: (70000, 784)
Shape of y: (70000,)

```
In [4]: plt.figure(figsize=(10, 5)) # setting size of pic

for i in range(40): # view first 40 pic
    plt.subplot(4, 10, i+1)
    plt.imshow(X[i].reshape(28, 28), cmap='gray')
    plt.title('Label: ' + str(y[i]))
    plt.axis('off')

plt.tight_layout()
plt.show()
```



```
In [5]: # View basic statistics of your data #查看数据类型
print("Data type:", X.dtype)
```

Data type: float64

```
In [6]: # View basic statistics of your data #查看数据结构
print(pd.DataFrame(X).describe())
```

	0	1	2	3	4	5	6	7	\
count	70000.0	70000.0	70000.0	70000.0	70000.0	70000.0	70000.0	70000.0	
mean	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
std	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
min	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
25%	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
50%	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
75%	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
max	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

	8	9	...	774	775	776	\
count	70000.0	70000.0	...	70000.000000	70000.000000	70000.000000	
mean	0.0	0.0	...	0.197414	0.099543	0.046629	
std	0.0	0.0	...	5.991206	4.256304	2.783732	
min	0.0	0.0	...	0.000000	0.000000	0.000000	
25%	0.0	0.0	...	0.000000	0.000000	0.000000	
50%	0.0	0.0	...	0.000000	0.000000	0.000000	
75%	0.0	0.0	...	0.000000	0.000000	0.000000	
max	0.0	0.0	...	254.000000	254.000000	253.000000	

	777	778	779	780	781	782	\
count	70000.000000	70000.000000	70000.000000	70000.0	70000.0	70000.0	
mean	0.016614	0.012957	0.001714	0.0	0.0	0.0	
std	1.561822	1.553796	0.320889	0.0	0.0	0.0	
min	0.000000	0.000000	0.000000	0.0	0.0	0.0	
25%	0.000000	0.000000	0.000000	0.0	0.0	0.0	
50%	0.000000	0.000000	0.000000	0.0	0.0	0.0	
75%	0.000000	0.000000	0.000000	0.0	0.0	0.0	
max	253.000000	254.000000	62.000000	0.0	0.0	0.0	

	783
count	70000.0
mean	0.0
std	0.0
min	0.0
25%	0.0
50%	0.0
75%	0.0
max	0.0

[8 rows x 784 columns]

```
In [7]: # 导入必要的库
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
```

```
In [8]: # data pre-processing
# Preprocess the data
print("Preprocessing data...")
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
print("Data preprocessing completed.")

# split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_

print("Number of samples in training set:", X_train.shape[0])
```

```
print("Number of samples in test set:", X_test.shape[0])
print("Number of features:", X_train.shape[1])
```

Preprocessing data...

Data preprocessing completed.

Number of samples in training set: 56000

Number of samples in test set: 14000

Number of features: 784

```
In [9]: # standardized data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Training logistic regression model
print("Training logistic regression model...")
logistic_regression = LogisticRegression(max_iter=1000, random_state=42)
logistic_regression.fit(X_train_scaled, y_train)

# print logistic regression model
train_accuracy_lr = logistic_regression.score(X_train_scaled, y_train)
test_accuracy_lr = logistic_regression.score(X_test_scaled, y_test)
print("Accuracy of logistic regression on training set:", train_accuracy_lr)
print("Accuracy of logistic regression on test set:", test_accuracy_lr)
```

Training logistic regression model...

Accuracy of logistic regression on training set: 0.945625

Accuracy of logistic regression on test set: 0.9164285714285715

```
In [10]: # Train SVM model
print("Training SVM model...")
svm_model = SVC()
svm_model.fit(X_train, y_train)
print("SVM model training completed.")

# Evaluate SVM model
svm_train_accuracy = accuracy_score(y_train, svm_model.predict(X_train))
svm_test_accuracy = accuracy_score(y_test, svm_model.predict(X_test))
print("SVM model train accuracy:", svm_train_accuracy)
print("SVM model test accuracy:", svm_test_accuracy)
```

Training SVM model...

SVM model training completed.

SVM model train accuracy: 0.9899464285714286

SVM model test accuracy: 0.9764285714285714

```
In [11]: from sklearn.ensemble import RandomForestClassifier
# Train RandomForest model
print("Training random forest model...")
random_forest = RandomForestClassifier(n_estimators=100, random_state=42)
random_forest.fit(X_train, y_train)

# Evaluate RandomForest model
train_accuracy_rf = random_forest.score(X_train, y_train)
test_accuracy_rf = random_forest.score(X_test, y_test)
print("Accuracy of random forest on training set:", train_accuracy_rf)
print("Accuracy of random forest on test set:", test_accuracy_rf)
```

Training random forest model...

Accuracy of random forest on training set: 1.0

Accuracy of random forest on test set: 0.9672857142857143

In [16]: `import joblib`

`# Save models`

```
joblib.dump(logistic_regression, 'C:/Users/armen/Downloads/EC-utbildning/2024-V.  
joblib.dump(random_forest, 'C:/Users/armen/Downloads/EC-utbildning/2024-V.7-Mach  
joblib.dump(svm_model, 'C:/Users/armen/Downloads/EC-utbildning/2024-V.7-Machine  
print("Models saved successfully.")
```

Models saved successfully.

In []: